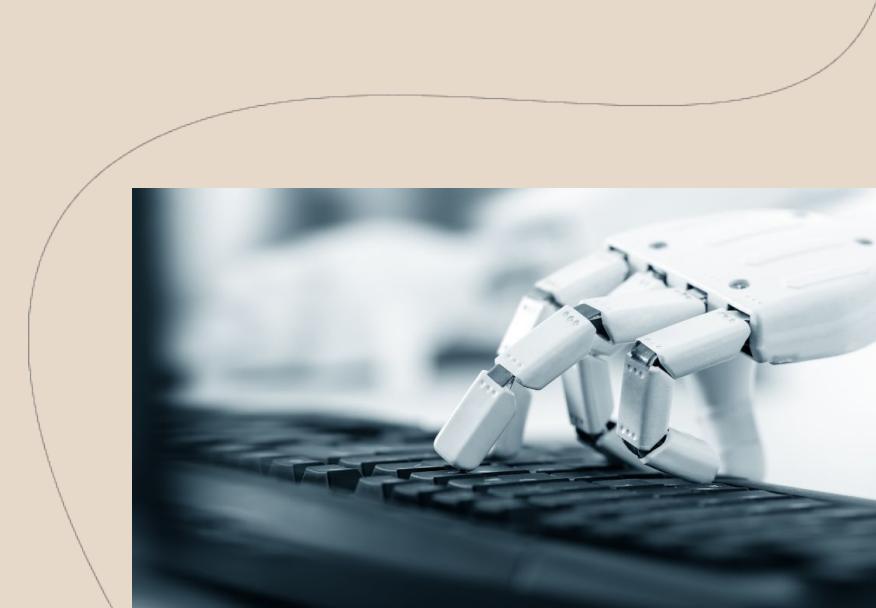




Postgraduaat

# AI Technical Architect

academiejaar '23-'24



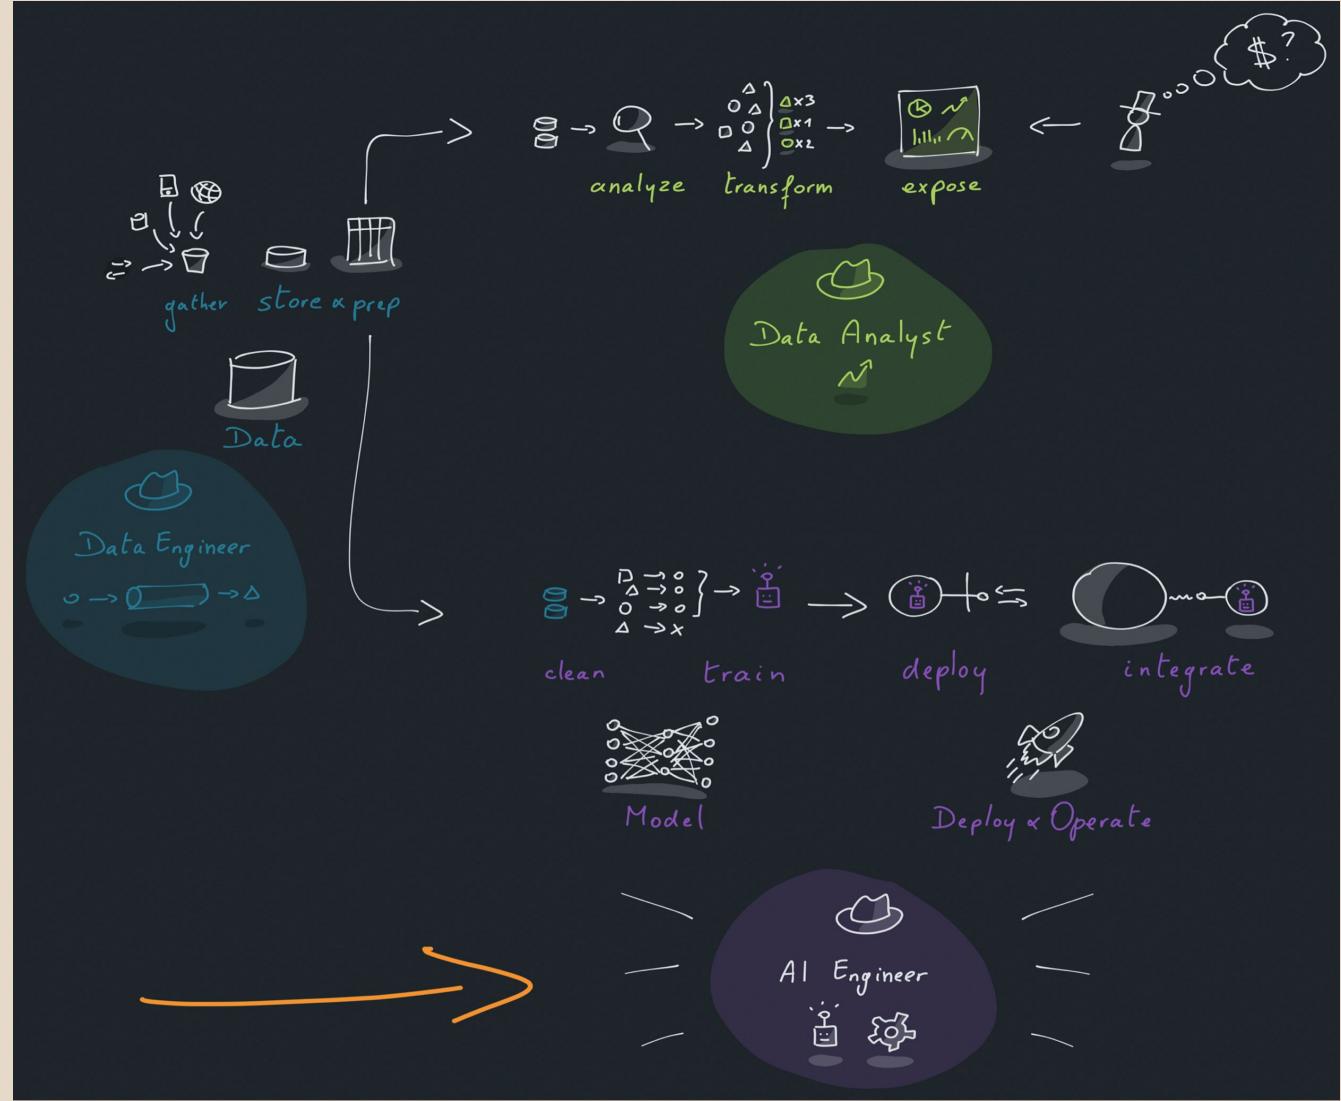


## Technologies *Hands on day 1*

*Tim Dupont*  
*Sam Van Rijn*

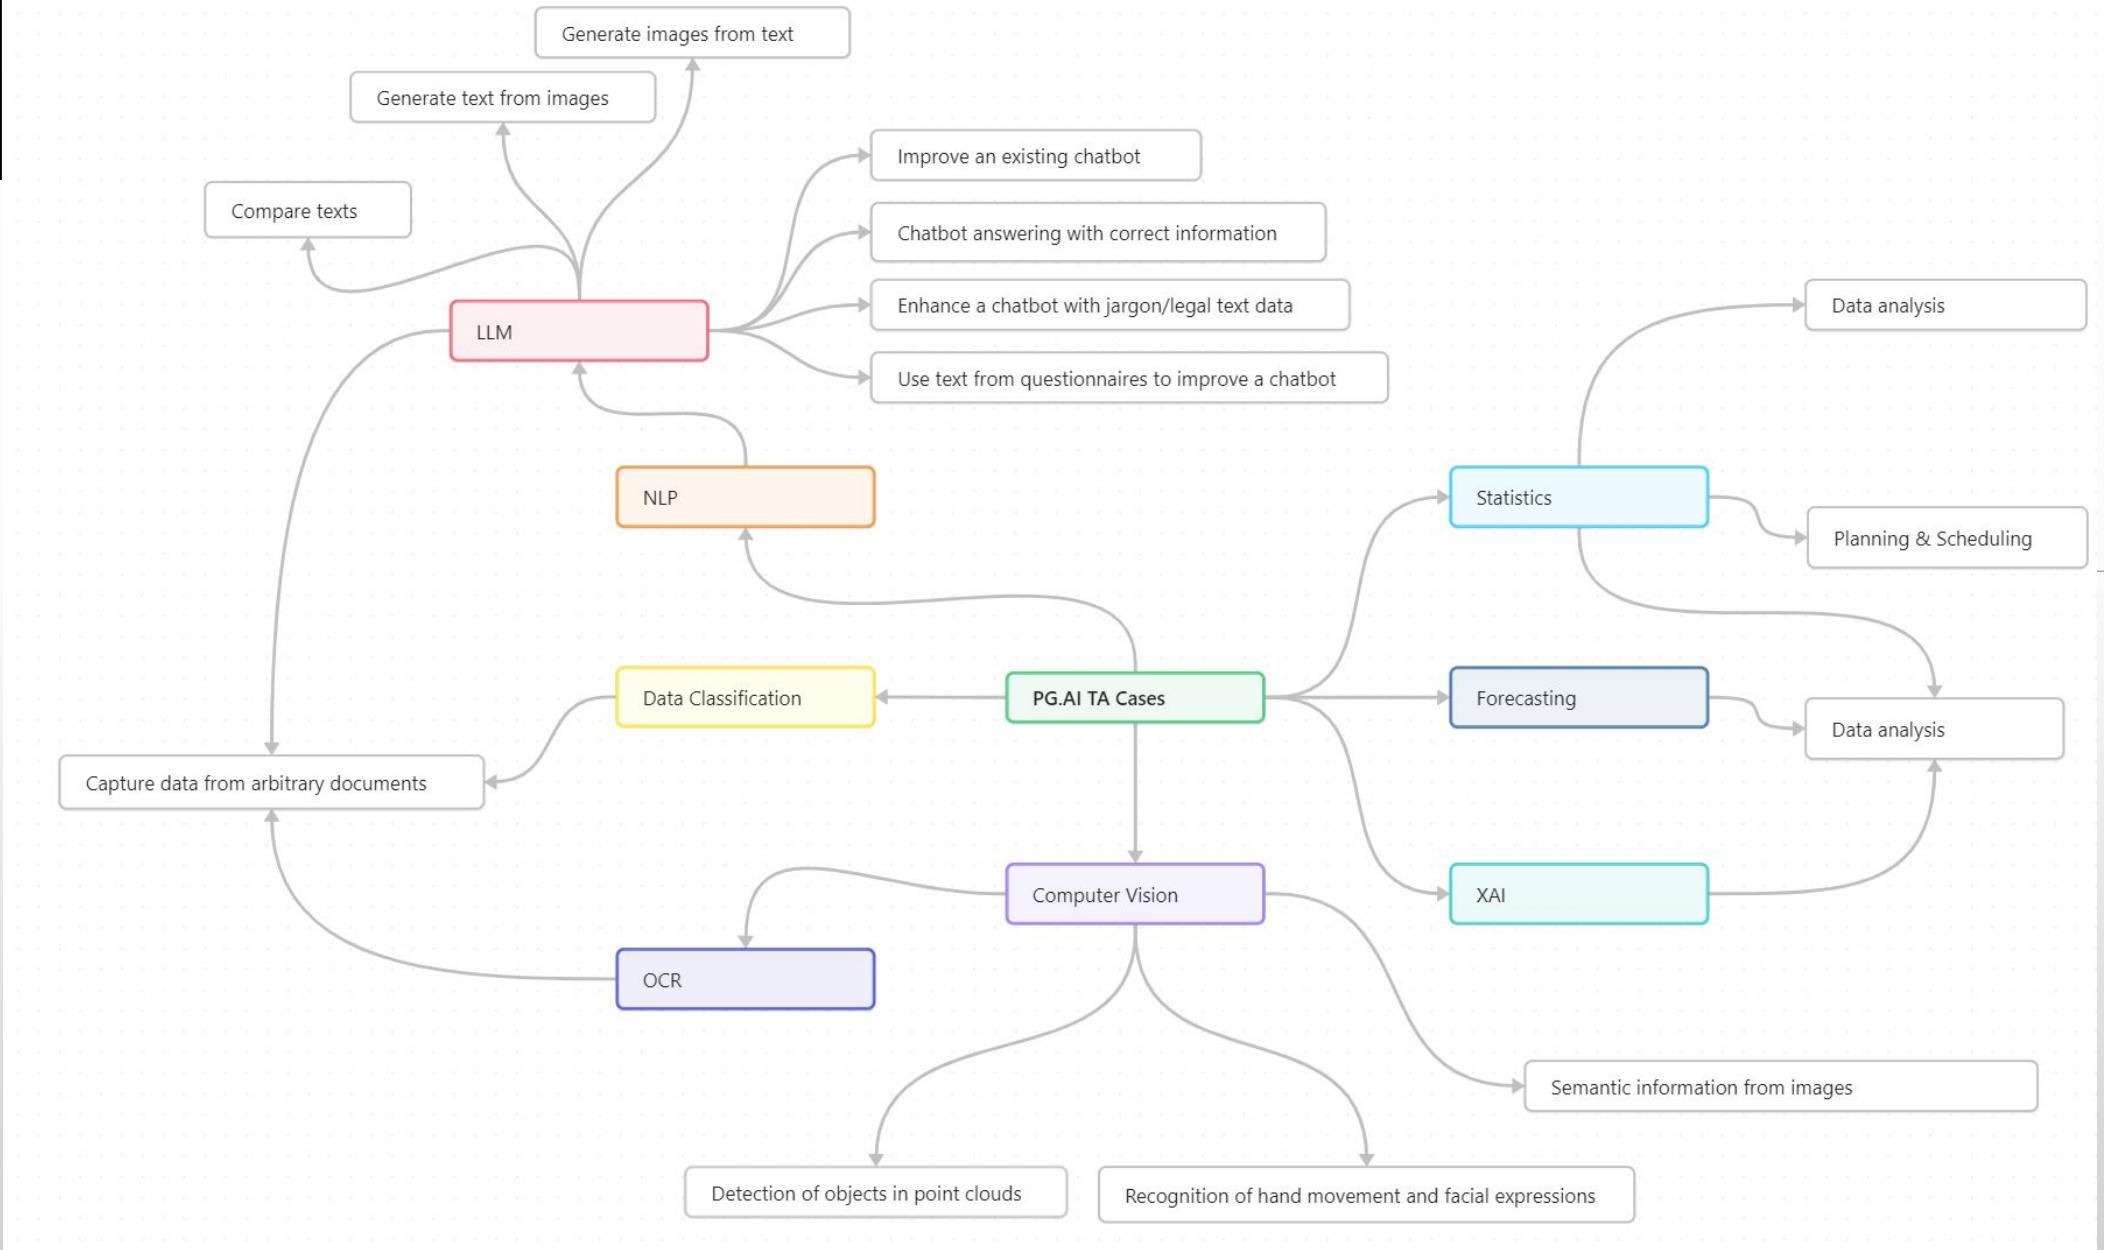
08 jan '24

# AI Technical Architect





N e X T



# Table of contents

1. Frameworks & libraries
  - a. Data Exploration & Model Training
    - i. Classical Machine Learning
    - ii. Deep Learning
  - b. Visualizations
  - c. Prototyping & Robotics
2. The tools we are going to use for the hands-on
3. Hands-on

## 01: First steps

- First Neural Network with Keras
- Credit Card Fraud

## 02: Optimizations via MNIST

## 03: CNNs & Transfer Learning

- Object detection, identification and transfer learning with YOLO
- CNN classifier

## 04: NLP (Next week)



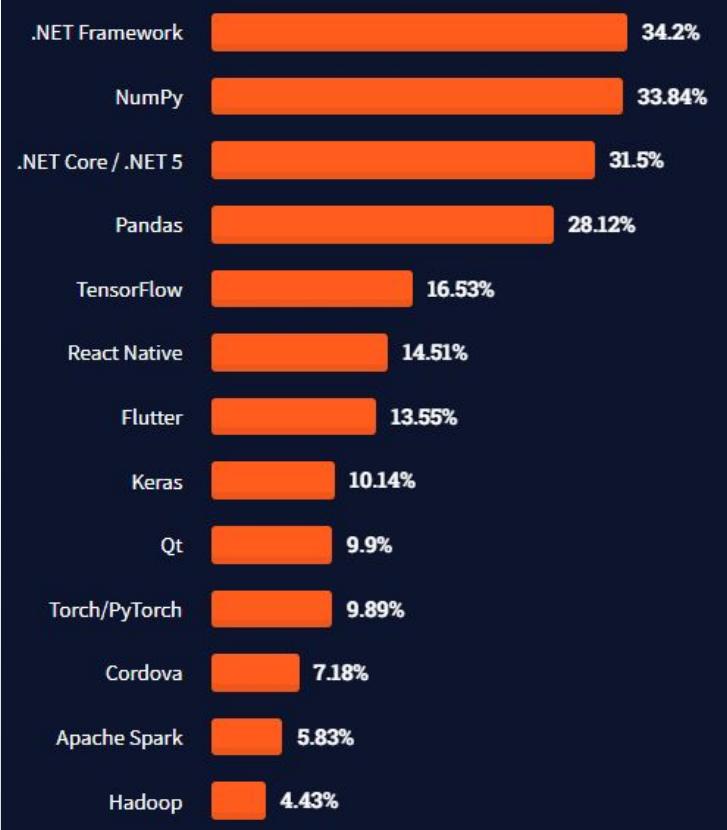
---

N e x T

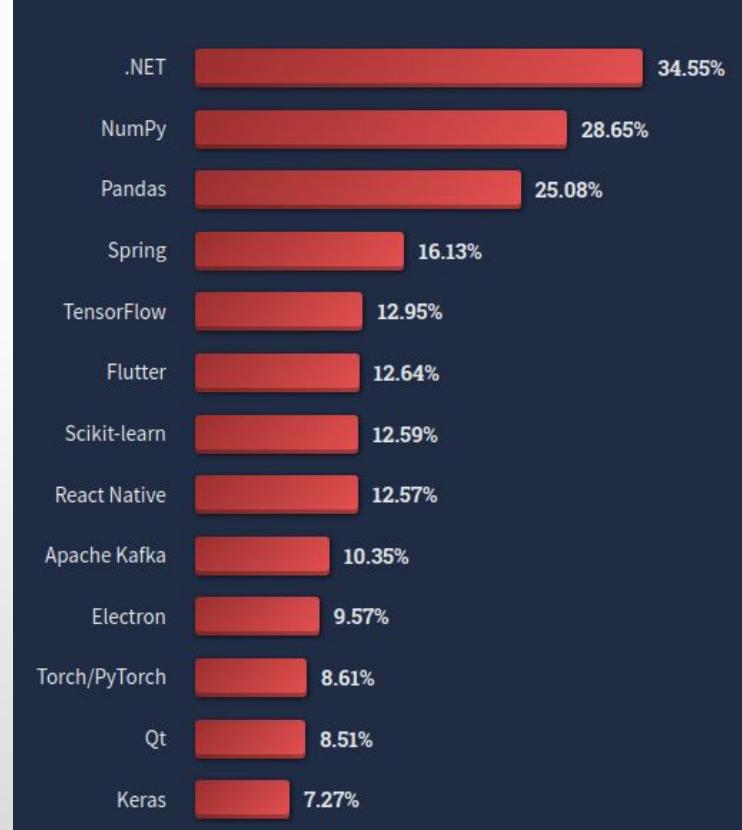
# Frameworks & libraries

# Most popular frameworks, libraries & tools

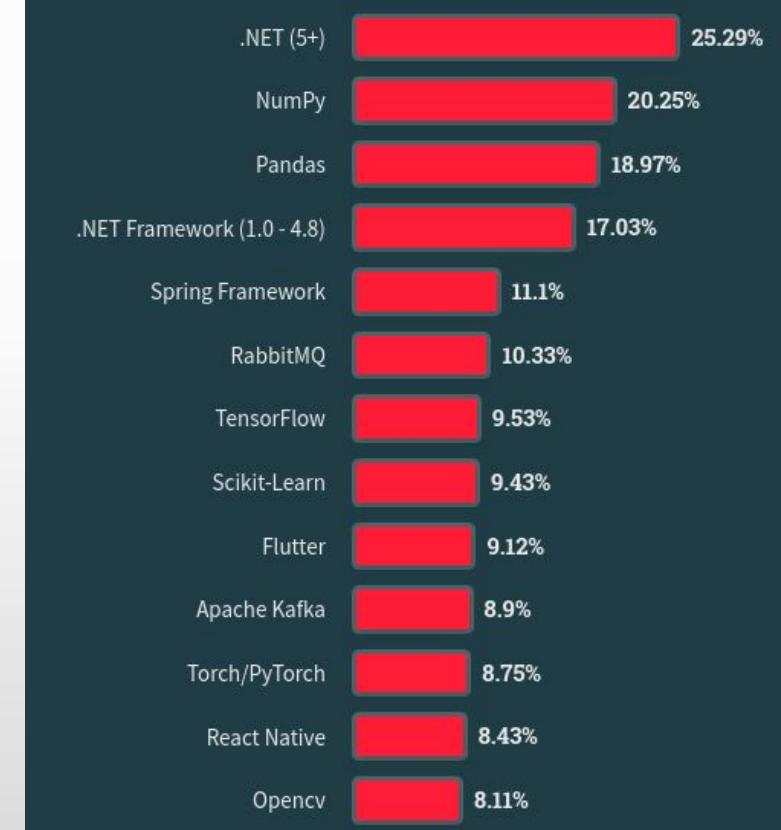
2021



2022



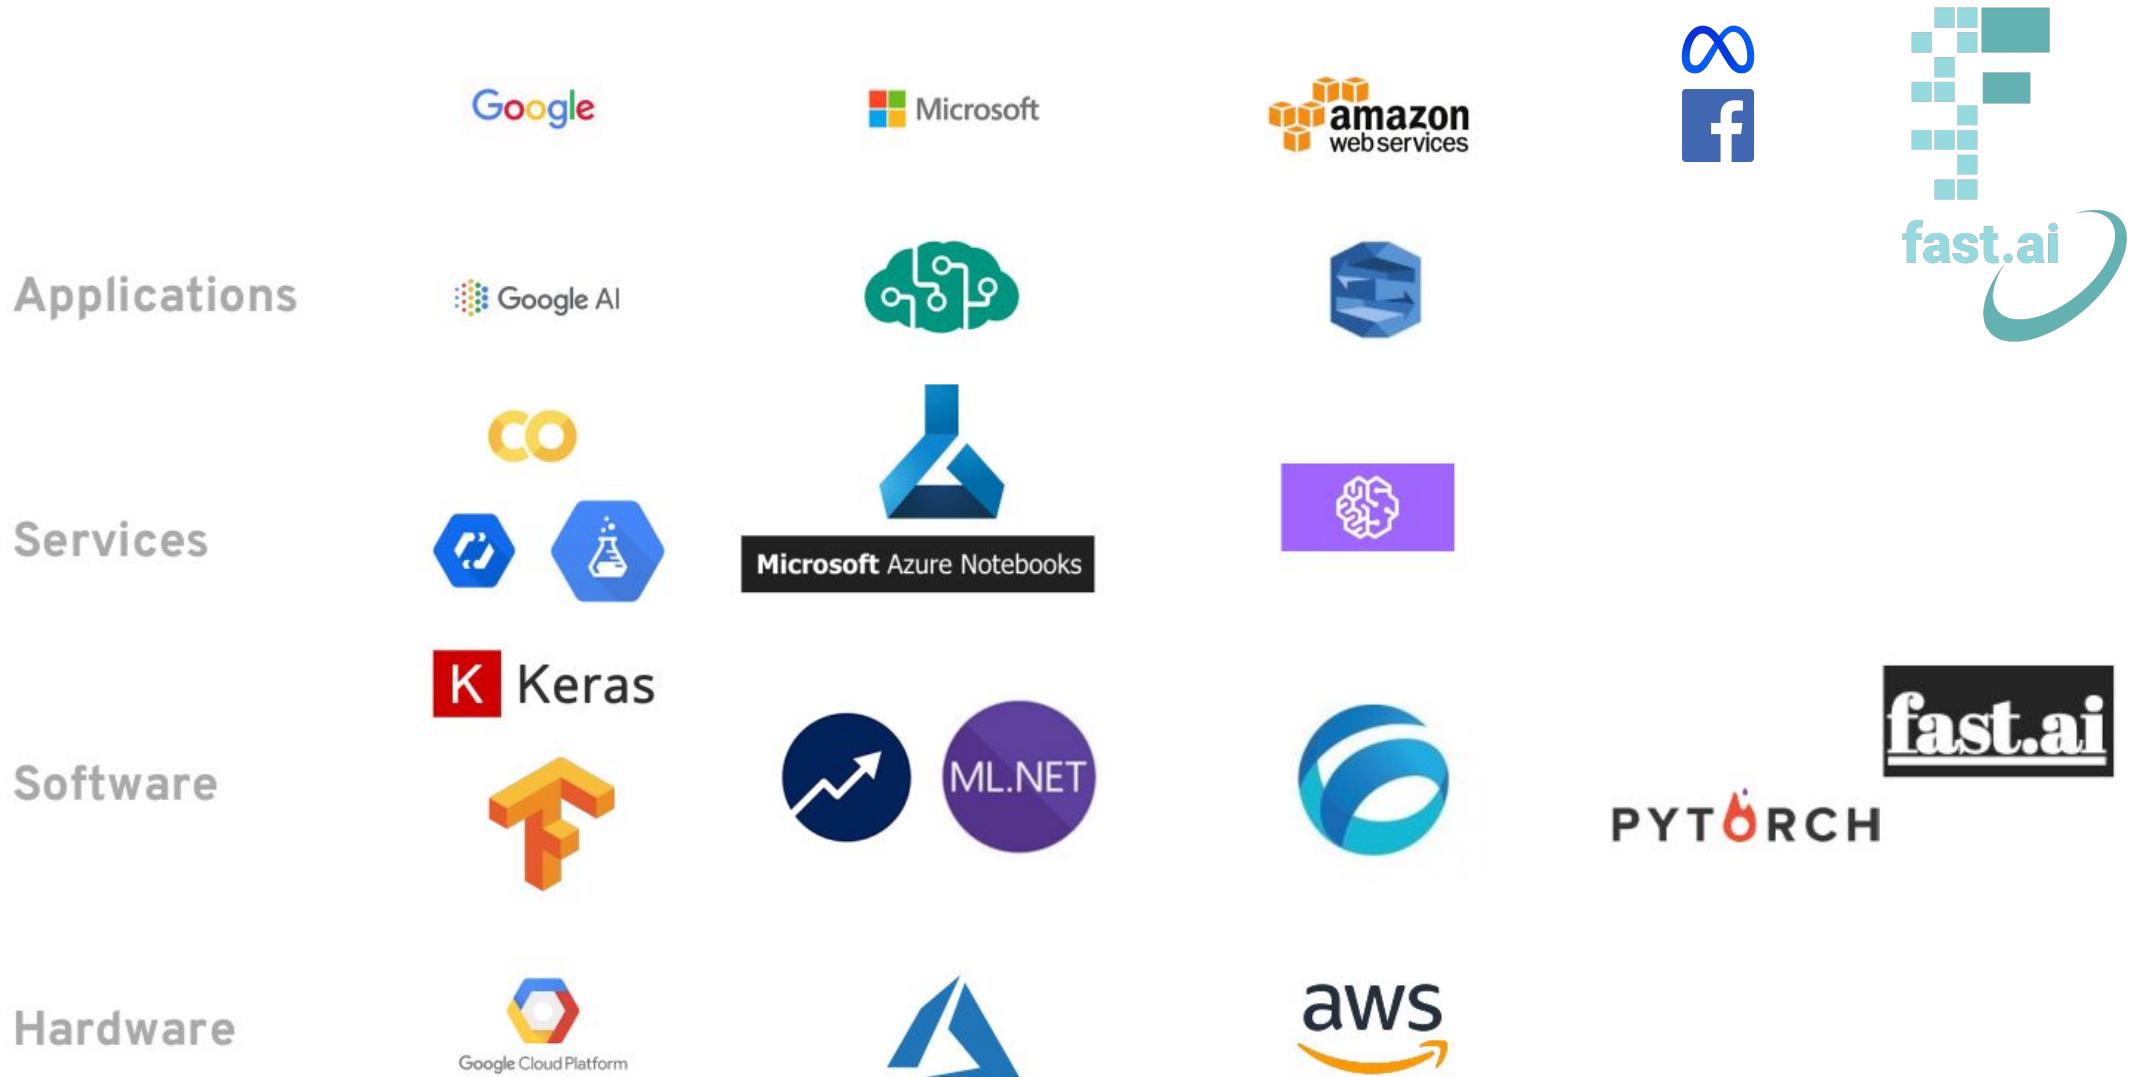
2023



[INFO]

Source: <https://insights.stackoverflow.com/survey/2021> & <https://survey.stackoverflow.co/2022> & <https://survey.stackoverflow.co/2023>

# AI software stack



# Python - Usage

Web development



Education

Prototyping



(Smart) Robotics



Scientific and Numeric



Business Applications



Cloud & Automation





---

N e x T

# Frameworks & libraries

## for Data Exploration & Model Training (classical Machine Learning)

# Data & ML



**SciPy**



# NumPy



- The fundamental package for scientific computing with Python
- Written in C and Python  
(In other words: “behind the scenes” is optimized, pre-compiled C code)
- Multidimensional array object
- Assortment of routines for fast operations on arrays  
(incl. manipulations, sorting, I/O, linear algebra, statistical operations, ...)
- Extensively used

Quantum Computing	Statistical Computing	Signal Processing	Image Processing	Graphs and Networks	Astronomy Processes	Cognitive Psychology
						
QuTiP PyQuil Qiskit	Pandas statsmodels Xarray Seaborn	SciPy PyWavelets python-control	Scikit-image OpenCV Mahotas	NetworkX graph-tool igraph PyGSP	AstroPy SunPy SpacePy	PsychoPy
Bioinformatics	Bayesian Inference	Mathematical Analysis	Chemistry	Geoscience	Geographic Processing	Architecture & Engineering
						
BioPython Scikit-Bio PyEnsembl ETE	PyStan PyMC3 ArviZ emcee	SciPy SymPy cvxpy FEniCS	Cantera MDAnalysis RDKit	Pangeo Simpeg ObsPy Fatiando a Terra	Shapely GeoPandas Folium	COMPAS City Energy Analyst Sverchok

[INFO]

More info: <https://numpy.org>

# pandas



- Data analysis and manipulation tool (library)
- Written in Python, C (and of course it uses NumPy)
- Fast, powerful, flexible and easy to use
  - DataFrame object for data manipulation
  - Reading and writing data between in-memory data structures and different file formats.
  - Data alignment and integrated handling of missing data.
  - Reshaping and pivoting of data sets.
  - Label-based slicing, fancy indexing, and sub-setting of large data sets.
  - Column insertion & deletion and data set merging & joining
  - Time series-functionality

```
df = df.dropna()  
  
df['Age'] = df['Age'].fillna(df['Age'].median())  
  
df['Embarked'] = df['Embarked'].fillna(df['Embarked'].mode()[0])
```

```
for column in obj_columns:  
    df[column] = df[column].astype('category')
```

```
df = pd.get_dummies(df, columns=['Embarked', 'Title'])
```

[INFO]

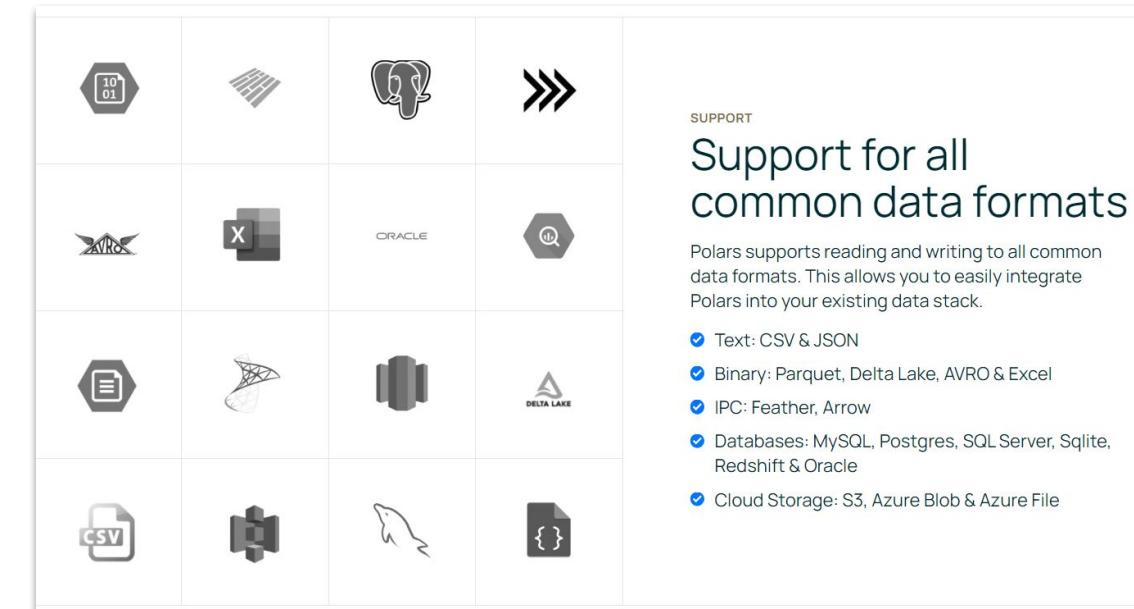
More info: <https://pandas.pydata.org>

# Polars



- Blazingly fast DataFrames in Rust, Python, Node.js, R and SQL
- Written in Rust
- Blazingly fast, lightweight...
  - Lazy | eager execution
  - Multi-threaded
  - SIMD
  - Query optimization
  - Powerful expression API
  - Hybrid Streaming (larger than RAM datasets)
  - Rust | Python | NodeJS | R | ...

```
import polars as pl\n\ndf = pl.read_csv("wowah_data.csv", parse_dates=False)\n\ndf.head(5)
```



[INFO]

More info: <https://pola.rs>

# SciPy



- Python library used for scientific computing and technical computing
- Written in Python, Fortran, C, C++ (and of course it uses NumPy)
- Cross platform
- Modules for
  - Optimization
  - Linear algebra
  - Integration
  - Interpolation
  - Fast Fourier Transform (FFT)
  - Signal and image processing
  - Clustering
  - ...

```
import scipy
from scipy.cluster import hierarchy as hc
corr = np.round(scipy.stats.spearmanr(X_keep).correlation, 4)
corr_condensed = hc.distance.squareform(1-abs(corr))
z = hc.linkage(corr_condensed, method='average')
fig = plt.figure(figsize=(16,10))
dendrogram = hc.dendrogram(z, labels=X_keep.columns,
                           orientation='left', leaf_font_size=16)
plt.show()
```

# scikit-learn



- Written in Python, C and C++ (and of course it uses NumPy)
- Supports most of the classical supervised and unsupervised learning algorithms:
  - Random forests
  - SVM
  - Naive Bayes
  - Gradient boosting
  - Clustering
- Designed to interoperate with the numerical and scientific libraries NumPy and SciPy
- Shipped with Anaconda

Advantages	Disadvantages
Great classical machine learning support	Very limited support for Neural Networks
Easy to use and excellent for beginners	No Deep Learning support

[INFO]

More info: <https://scikit-learn.org/stable/> and <https://github.com/scikit-learn/scikit-learn>

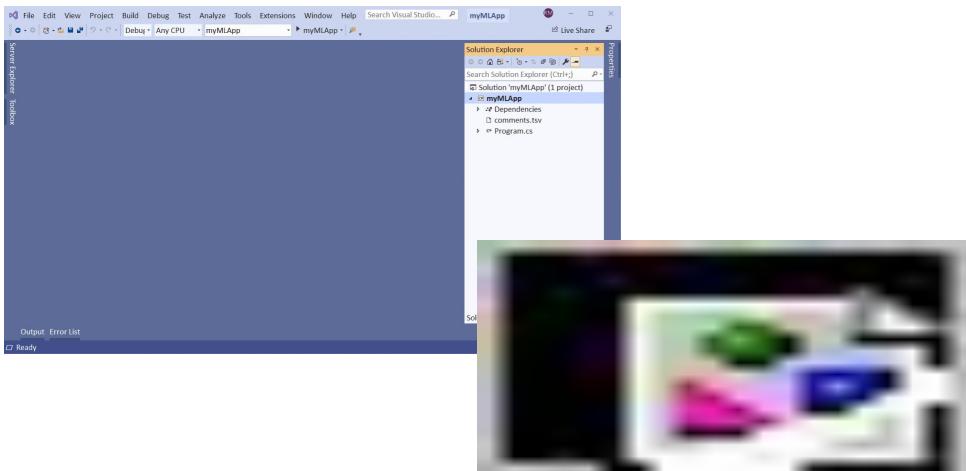
# ML.NET



- Model-based Machine Learning analytic and prediction capabilities
- For existing .NET developers (C# and F#)
- Written in C# and C++
- Built upon .NET Core and .NET Standard
- Cross platform (Linux, macOS, and Windows)
- Model Builder (a simple UI tool) and ML.NET CLI
- Documentation used to be awful, but it's getting better

(Intro: <https://learn.microsoft.com/en-us/dotnet/machine-learning/how-does-mldotnet-work>)

(Samples: <https://github.com/dotnet/machinelearning-samples>)



 <b>Sentiment analysis</b> Analyze the sentiment of customer reviews using a binary classification algorithm.	 <b>Product recommendation</b> Recommend products based on purchase history using a matrix factorization algorithm.	 <b>Price prediction</b> Predict taxi fares based on parameters such as distance traveled using a regression algorithm.
 <b>Customer segmentation</b> Identify groups of customers with similar profiles using a clustering algorithm.	 <b>Object detection</b> Recognize objects in an image using an ONNX deep learning model.	 <b>Fraud detection</b> Detect fraudulent credit card transactions using a binary classification algorithm.
 <b>Sales spike detection</b> Detect spikes and changes in product sales using an anomaly detection model.	 <b>Image classification</b> Classify images (for example, broccoli vs. pizza) using a TensorFlow deep learning model.	 <b>Sales forecasting</b> Forecast future sales for products using a regression algorithm.

You can find [more ML.NET samples on GitHub](#), or take a look at the [ML.NET tutorials](#).

## [INFO]

More info: <https://dotnet.microsoft.com/en-us/apps/machinelearning-ai/ml-dotnet> and <https://github.com/dotnet/machinelearning>



---

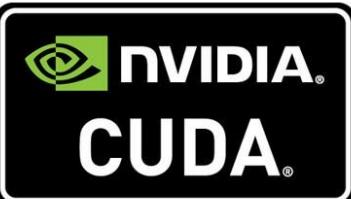
N e x T

# Frameworks & libraries

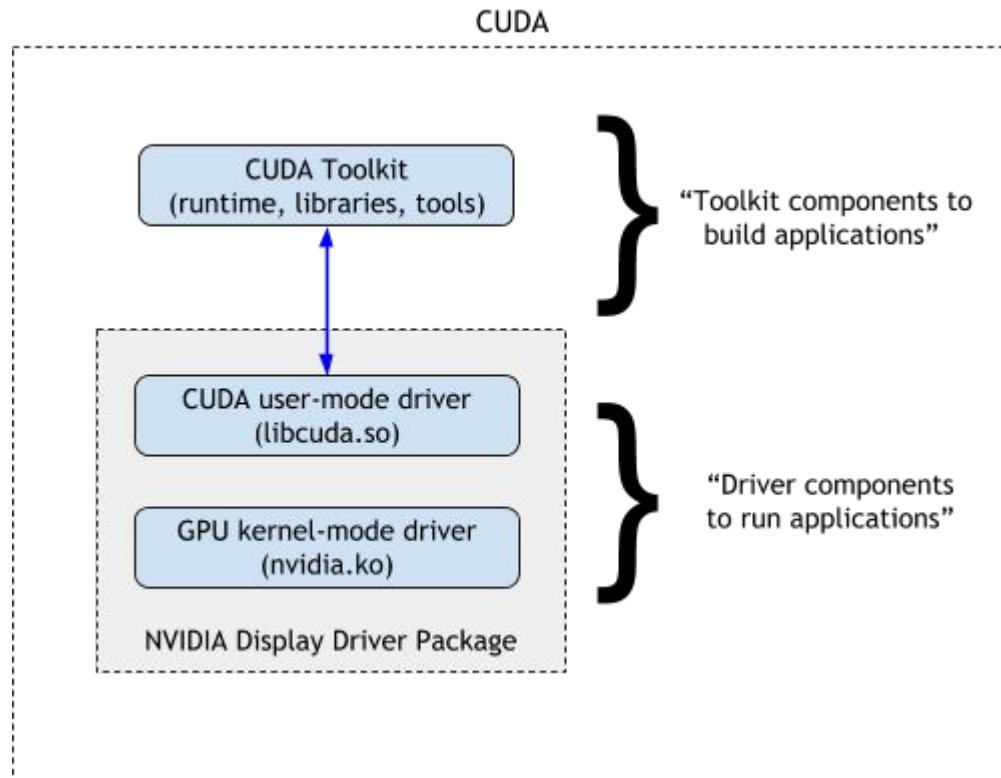
## for Data Exploration & Model Training (Deep Learning)

# Deep learning





# Nvidia CUDA

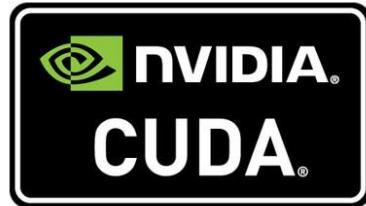


- CUDA (Compute Unified Device Architecture)
- Parallel computing platform + Application Programming Interface (API)
- Use GPU for general purpose processing (GPU has parallel processors called CUDA Cores  
example: GeForce RTX 3090 has 10496 cores)
- libcudnn: C Lib for Deep learning (Deep Neural Network library)
- **By far the most popular AI computing platform!**  
**Supported by most AI libraries:**  
Keras, MATLAB, PyTorch, TensorFlow, ...
- **Only for Nvidia graphics cards!**  
(OpenCL is an open alternative, but less popular)

[INFO]

More info: <https://developer.nvidia.com/cuda-toolkit>

# Nvidia CUDA



## FRAMEWORKS



Chainer



mxnet



PaddlePaddle



PyTorch



TensorFlow

## CLOUD ML SERVICES



## DEPLOYMENT



DA

GRAPH

ML

DL TRAIN

DL INFERENCE

CUDA-X AI

CUDA

Workstation



Server



Cloud

[INFO]

More info: <https://www.nvidia.com/en-us/technologies/cuda-x>



# Nvidia CUDA

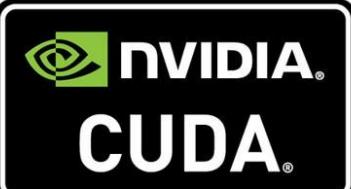
A screenshot of a terminal window titled "Ubuntu-22.04". The command "nvidia-smi" is run at the prompt. The output shows GPU information and no running processes.

```
tim@othronn:~$ nvidia-smi
Mon Jan  8 09:00:03 2024
+-----+
| NVIDIA-SMI 525.147.01   Driver Version: 529.19      CUDA Version: 12.0 |
+-----+
| GPU  Name Persistence-M| Bus-Id     Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf Pwr:Usage/Cap| Memory-Usage | GPU-Util Compute M. |
|                               |             |          MIG M. |
+-----+
| 0  NVIDIA RTX A200... On | 00000000:01:00.0 Off | N/A |
| N/A  52C    P0    10W / 35W | 0MiB / 4096MiB | 0% Default |
|                               |             |          N/A |
+-----+
+-----+
| Processes:
| GPU  GI CI PID Type Process name          GPU Memory |
| ID   ID          ID   ID                 Usage
+-----+
| No running processes found
+-----+
tim@othronn:~$
```

Checking GPU with `nvidia-smi` in WSL2 (Ubuntu-22.04 on Windows 11).

[INFO]

Check with: `nvidia-smi`



# Nvidia CUDA

A screenshot of a Google Colab notebook titled "CudaTestOnColab.ipynb". The code cell contains the command "!nvidia-smi". The output shows GPU information for a Tesla T4, including power usage, memory, and compute utilization. It also lists processes running on the GPU.

```
Mon Jan  8 08:49:11 2024
+-----+
| NVIDIA-SMI 535.104.05      Driver Version: 535.104.05    CUDA Version: 12.2 |
+-----+
| GPU  Name     Persistence-M | Bus-Id     Disp.A  Volatile Uncorr. ECC | | | | |
| Fan  Temp     Perf            Pwr:Usage/Cap | Memory-Usage | GPU-Util Compute M. |
|          |          |          |           |           |          MIG M. |
+-----+
| 0  Tesla T4          Off  00000000:00:04.0 Off   0% Default |
| N/A  35C   P8          9W / 70W |    0MiB / 15360MiB |       0%      N/A |
+-----+
+-----+
| Processes:                               GPU Memory |
| GPU  GI  CI      PID  Type  Process name        Usage  |
| ID   ID          ID          |
+-----+
| No running processes found               |
+-----+
```

Checking GPU with nvidia-msi on Google Colab

[INFO]

Check with: nvidia-msi

# TensorFlow



- Python
- Created by Google
- Written in Python, C++ (and of course it uses NumPy)
- Makes use of CUDA
- Tensor computation (like NumPy) with strong GPU acceleration
- Stateful dataflow graphs
- Tensor Processing Unit (TPU) → dedicated hardware support

Advantages	Disadvantages
Fast	Not so easy to use (too low level)
Neural networks and Deep Learning support	

[INFO]

More info: <https://www.tensorflow.org> and <https://github.com/tensorflow/tensorflow>

# Keras



- Python library for artificial neural networks
- On top of TensorFlow a.k.a. interface for the TensorFlow library  
(Until version 2.3 also support for Microsoft Cognitive Toolkit, Theano, or PlaidML.)
- Strong ties with Google

Advantages	Disadvantages
Machine Learning and Deep Learning support	Not enough documentation
Easy to use	

```
network = models.Sequential()
network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))

# The final layer has 10 nodes, each node represents one class of numbers
network.add(layers.Dense(10, activation='softmax'))

network = models.Sequential()
network.add(layers.Dense(300, activation='relu', input_shape=(28 * 28,), kernel_regularizer=keras.regularizers.l1_l2(l1=0, l2=0.001)))
#network.add(layers.Dropout(0.4))
network.add(layers.Dense(100, activation='relu'))#, kernel_regularizer=l2(0.1)))
#network.add(layers.Dense(300, activation='relu', kernel_regularizer=l2(0.0001)))
#network.add(layers.Dropout(0.2))

# The final layer has 10 nodes, each node represents one class of numbers
network.add(layers.Dense(10, activation='softmax'))
```

[INFO]

More info: <https://keras.io> and <https://github.com/keras-team/keras>

# PyTorch

- Written in Python, C++, CUDA (and of course it uses NumPy)
- Created by Facebook
- Tensor computation (like NumPy) with strong GPU acceleration
- Deep neural networks built on a tape-based autodiff system

Advantages	Disadvantages
Machine Learning and Deep Learning support	Not enough documentation about some methods and parameters
Allows customization	

[INFO]

More info: <https://pytorch.org> and <https://github.com/pytorch/pytorch>

- Written in Python (and of course it uses NumPy)
- Built on top of Scikit-learn and PyTorch
- Aim is to make Machine and Deep Learning as accessible as possible
- Contains a lot of functions to make data processing easier

Advantages	Disadvantages
Great all round support	API changes quite often
Integrates new techniques and methods rapidly	Less market usage
Great for learning concepts easily	



---

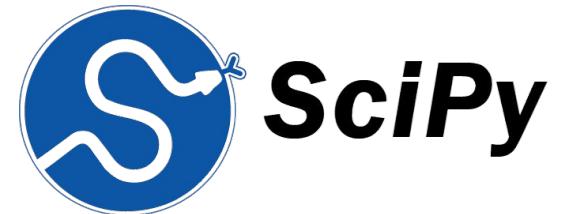
N e x T

# Frameworks & libraries for Visualizations

# Visualizations



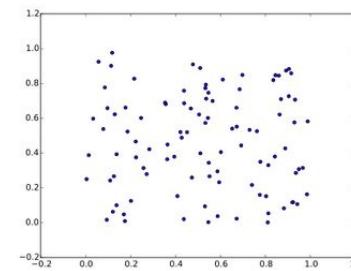
matplotlib



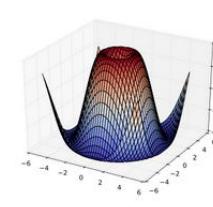
# Matplotlib



- Python library for creating static, animated, and interactive visualizations
- Written in Python
- Makes extensive use of NumPy
- Embeddable in GUI toolkits (i.e. QT)



Scatter plot



3D plot

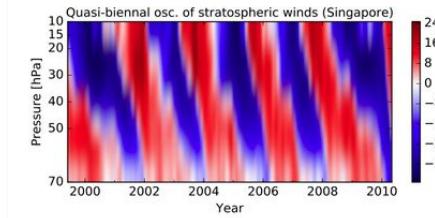
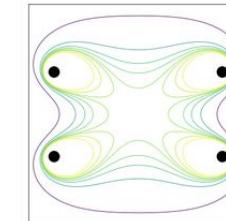
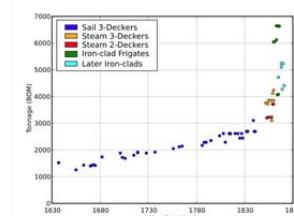


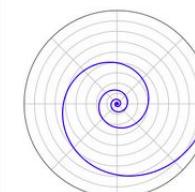
Image plot



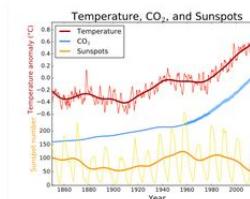
Contour plot



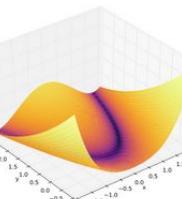
Scatter plot



Polar plot



Line plot



3-D plot

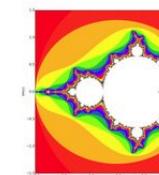
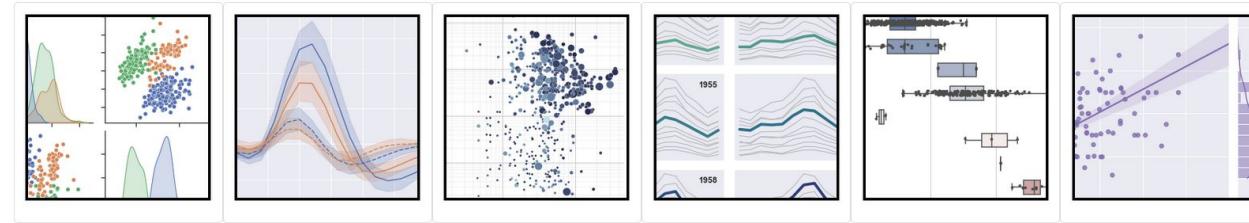


Image plot

[INFO]

More info: <https://matplotlib.org>

# Seaborn



- Data visualization library based on Matplotlib
- High-level interface for drawing attractive and informative statistical graphics
- Integrates closely with pandas data structures

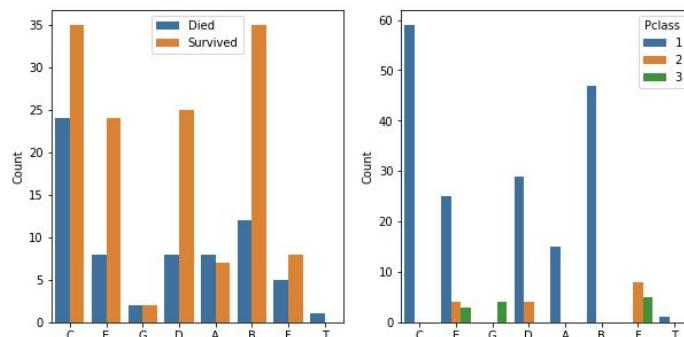
```
fig, ax = plt.subplots(1,2, figsize=(10,5))

sns.countplot('Deck', data=df[df['Deck'] != 'Z'], hue='Survived', ax=ax[0])
sns.countplot('Deck', data=df[df['Deck'] != 'Z'], hue='Pclass', ax=ax[1])

ax[0].set_xlabel('Deck')
ax[0].set_ylabel('Count')
ax[0].legend(['Died', 'Survived'])

ax[1].set_xlabel('Deck')
ax[1].set_ylabel('Count')

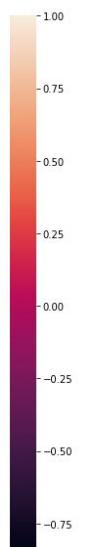
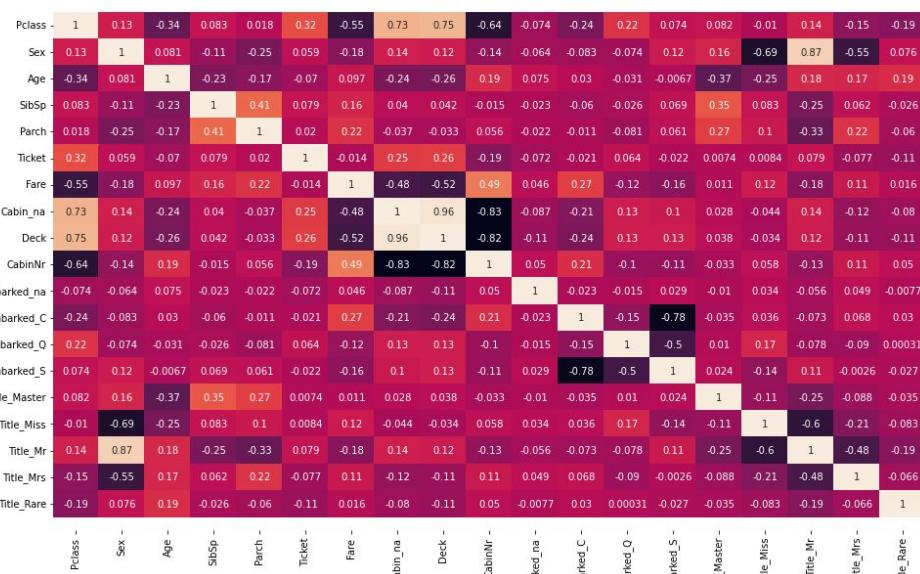
fig.show()
```



```
fig, ax = plt.subplots(figsize = (20, 10))
# we drop the feature 'Survived', since this is our dependent variable (the variable we are trying to predict)
df_corr = df.drop(['Survived'], axis=1).corr()
fig = sns.heatmap(df_corr, annot = True)

# This is a fix for a bug in the visualization library
bottom, top = ax.get_ylim()
ax.set_ylim(bottom + 0.5, top - 0.5)
```

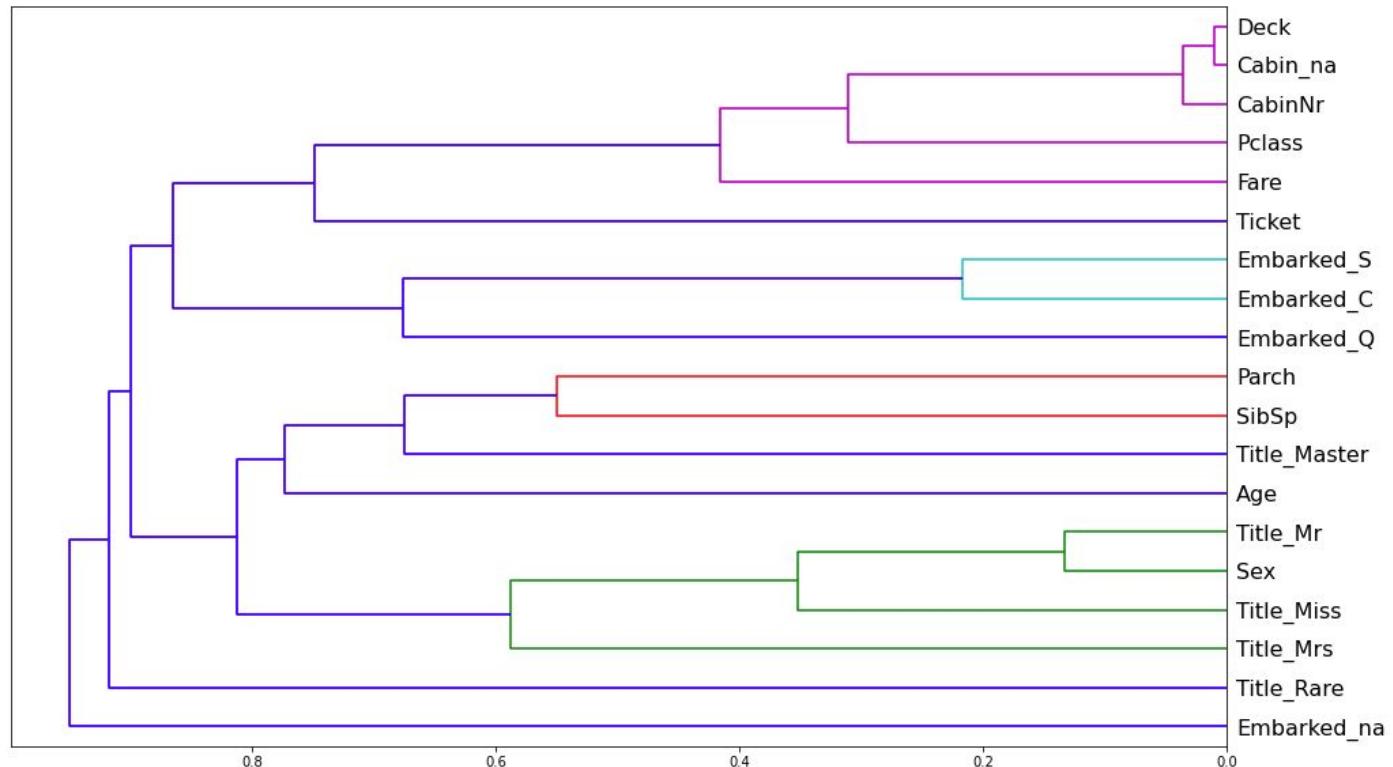
(19.5, -0.5)



[INFO]

More info: <https://seaborn.pydata.org/introduction.html>

```
import scipy
from scipy.cluster import hierarchy as hc
corr = np.round(scipy.stats.spearmanr(df.drop(['Survived'], axis=1)).correlation, 4)
corr_condensed = hc.distance.squareform(1-abs(corr))           # we use 1-abs(corr) instead of 1-corr, to incorporate negative correlations as well as positive ones
z = hc.linkage(corr_condensed, method='average')
fig = plt.figure(figsize=(16,10))
dendrogram = hc.dendrogram(z, labels=df.columns, orientation='left', leaf_font_size=16)
plt.show()
```



[INFO]

More info: <https://scipy.org> and <https://github.com/scipy/scipy>

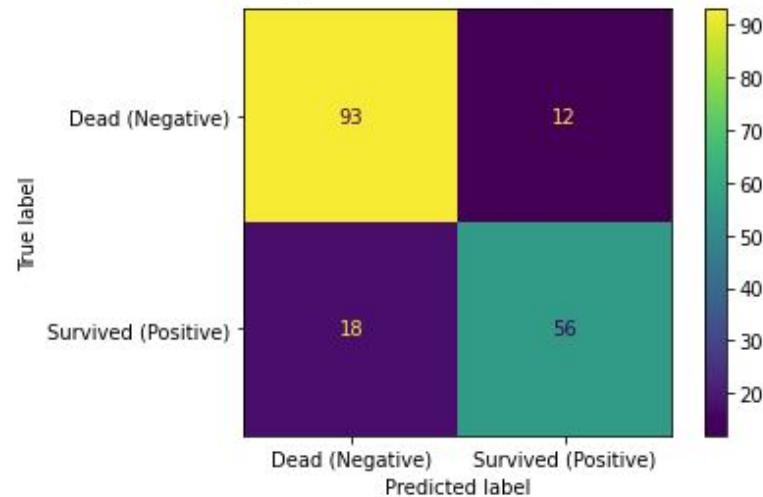
# scikit-learn



```
RF_preds = RF_model.predict(X_test)
RF_preds_rounded = np.round_(RF_preds, 0)
RF_cm = confusion_matrix(y_test, RF_preds_rounded)

disp = ConfusionMatrixDisplay(confusion_matrix=RF_cm,
                               display_labels=['Dead (Negative)',
                                              'Survived (Positive)'])
disp.plot()

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7fe34696f310>
```



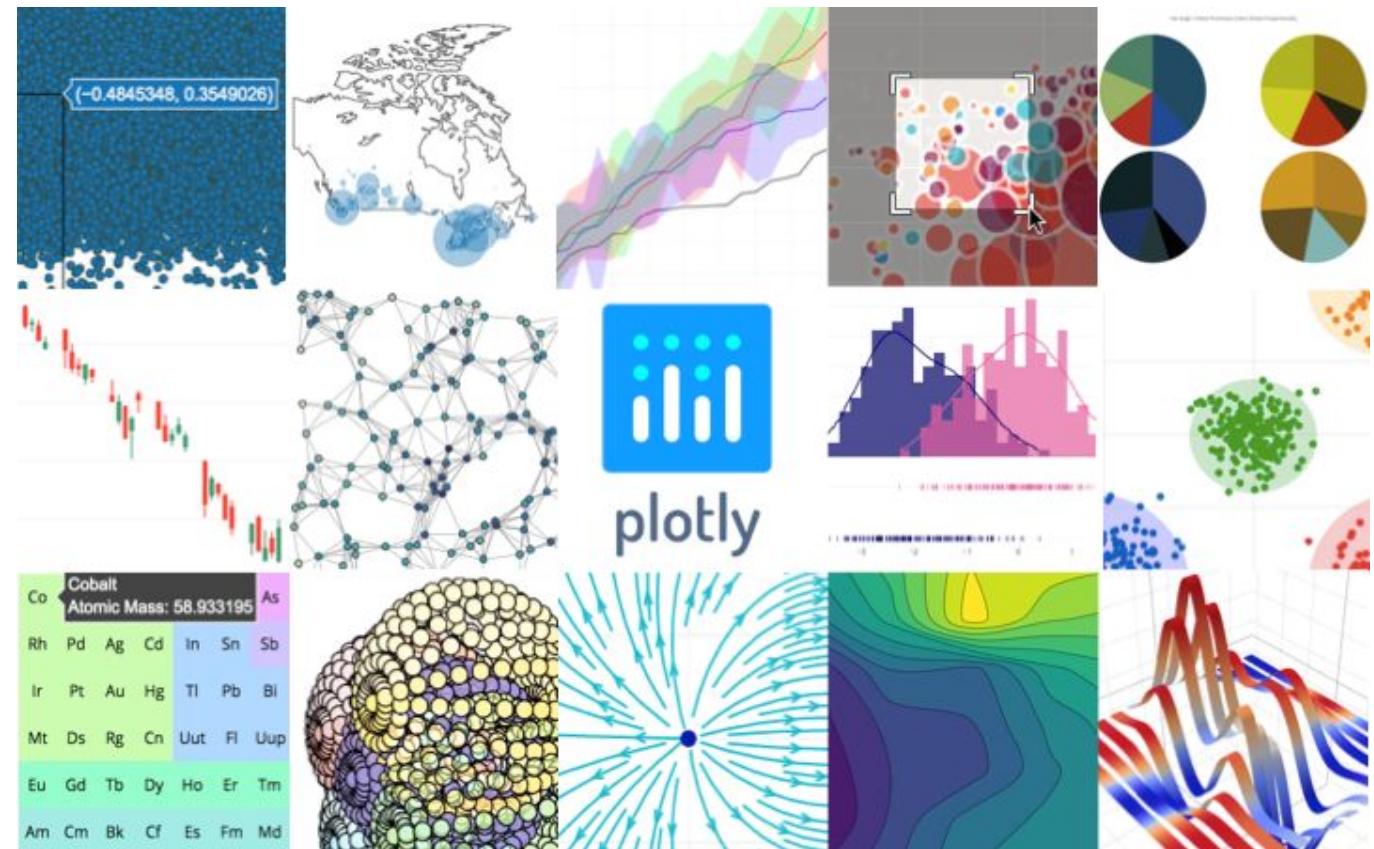
[INFO]

More info: <https://scikit-learn.org/stable/> and <https://github.com/scikit-learn/scikit-learn>

# plotly



- Interactive, open-source, and browser-based graphing library for Python
- Built on top of plotly.js,
- Declarative charting library.
- Over 30 chart types, including:
  - Scientific charts
  - 3D graphs
  - Statistical charts
  - SVG maps
  - Financial charts
  - ...



[INFO]

More info: <https://plotly.com/python>



# Data-Driven Documents (d3.js)

- JavaScript library for manipulating documents based on data.
- Using HTML, SVG, and CSS.
- Emphasis on web standards



[INFO]

More info: <https://d3js.org>

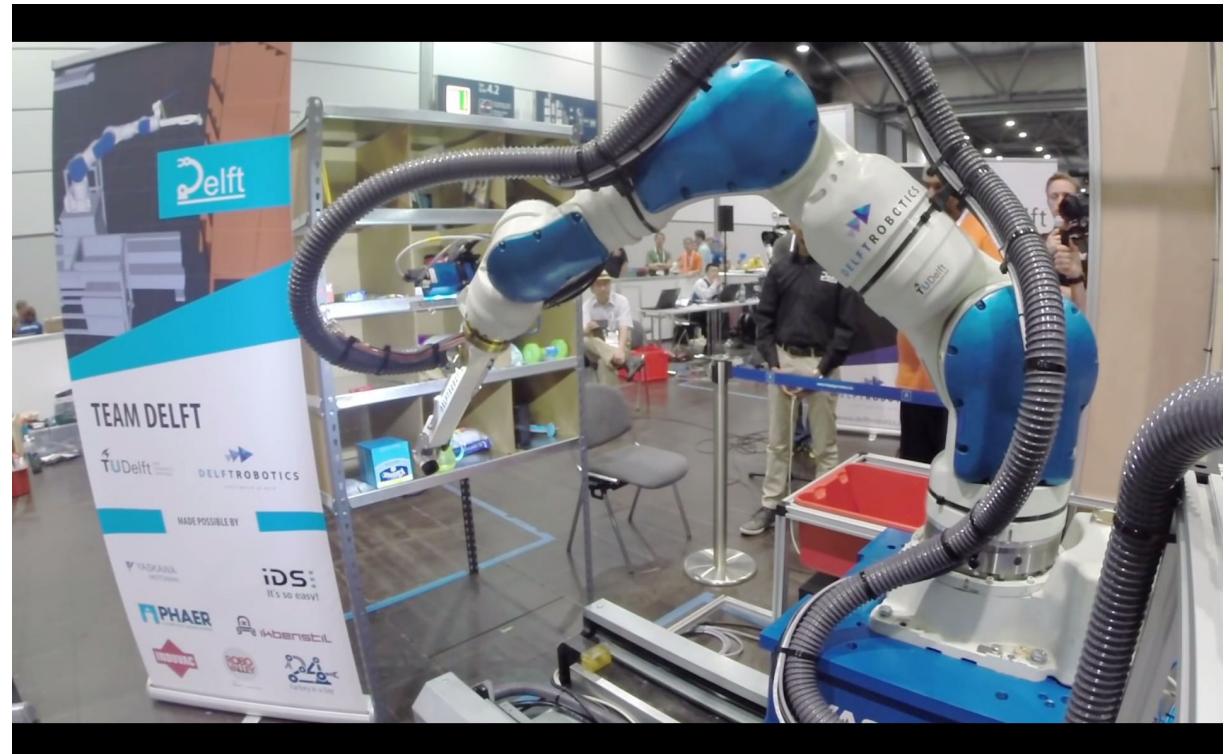
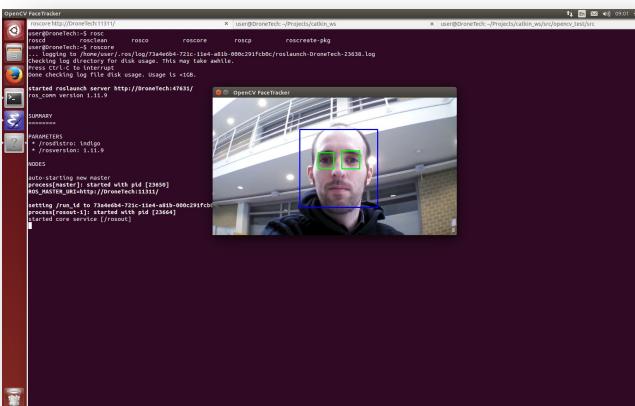
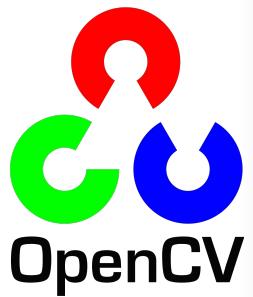


# Frameworks & libraries

## Prototyping & Robotics

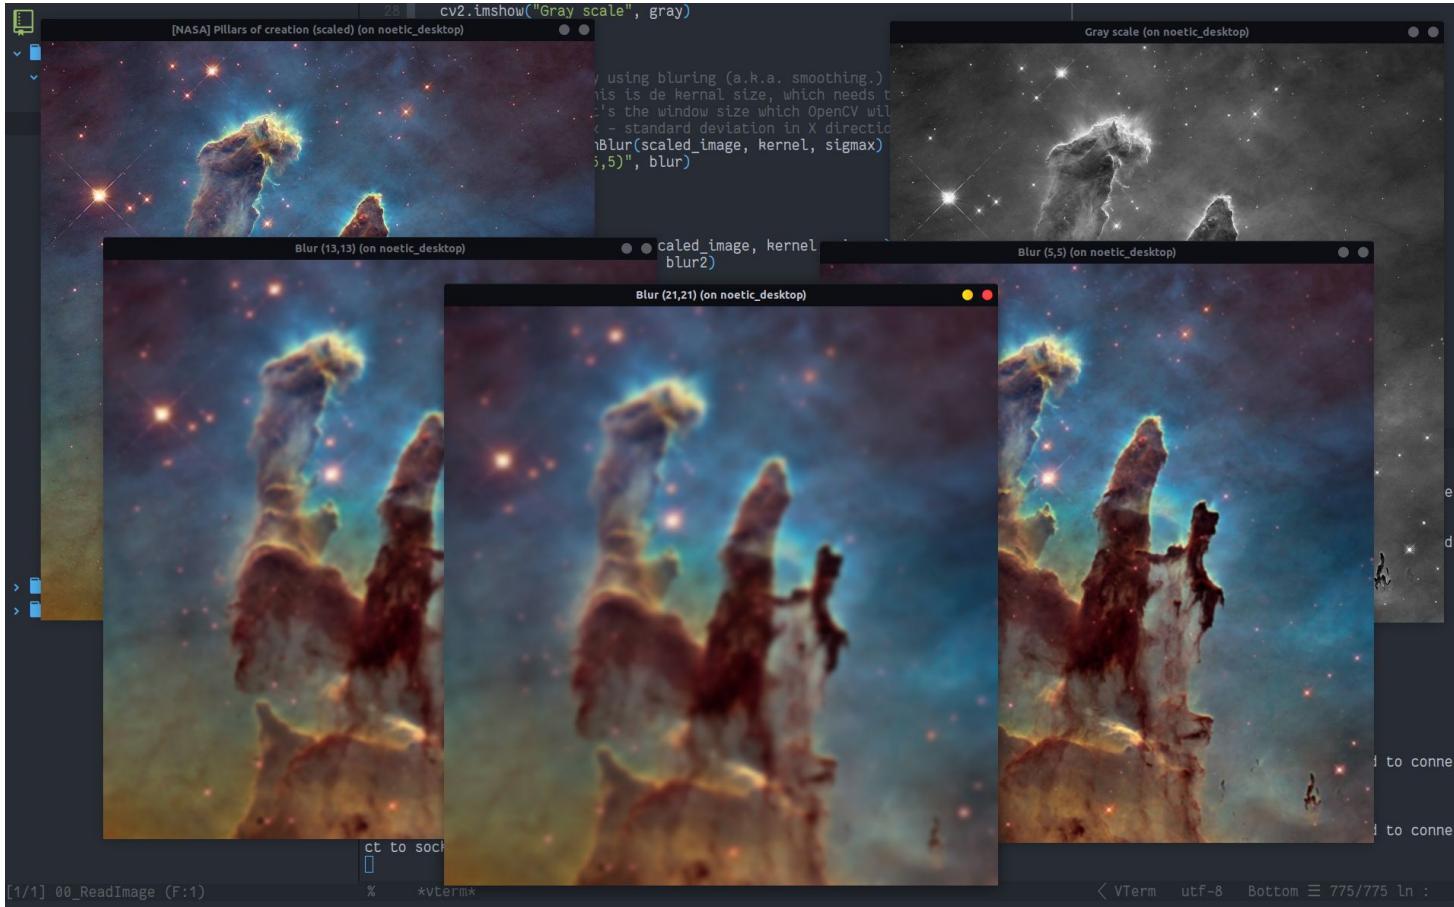
# Prototyping & Robotics

Not strictly Python





# OpenCV

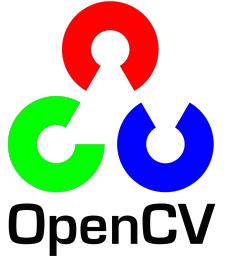


- Open Source Computer Vision
- Cross-platform library
- Real-time Computer Vision
- Originally by Intel
- Well documented
- Multiple bindings  
(C++, Java & Python)

[INFO]

More info: <https://opencv.org>

# OpenCV in 4h



# Python & OpenCV in 4 hours



()

SUBSCRIBE

[INFO]

Source: <https://www.youtube.com/watch?v=oXlwWbU8l2o>

# QT



Panasonic Avionics –  
Inflight Infotainment

[Read More](#)



Veriskin – Skin Cancer  
Diagnostics Device

[Read More](#)



Rimac Automobili – Car  
Infotainment System

[Read More](#)



AutoIO – Digital  
Instrument Clusters

[Read More](#)



BEP Marine – Marine  
Automation Systems

[Read More](#)



AMD – Graphics software

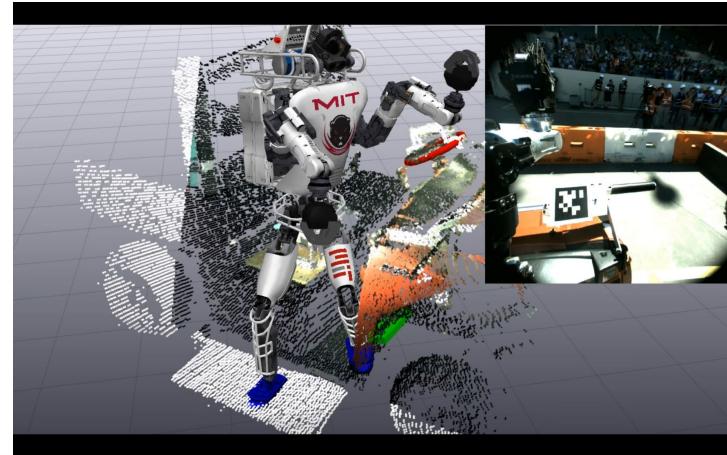
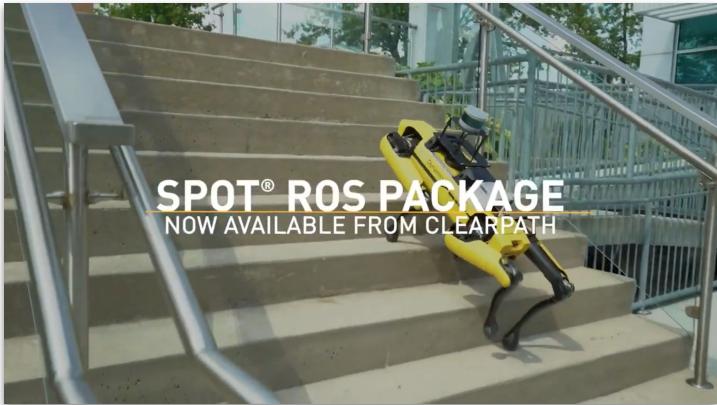
[Read More](#)

(pronounced "cute") is a **free and open-source widget toolkit** for creating graphical user interfaces as well as cross-platform applications that **run on various software and hardware platforms** such as Linux, Windows, macOS, Android or embedded systems with **little or no change in the underlying codebase** while still being a **native application with native capabilities and speed**.

[INFO]

More info: <https://www.qt.io>

# ROS



- Open source software (BSD licence)
- Robotics framework
- C++, Python, ...
- “Flexible Robotics”
- Message framework
- Active Community
- Used in industry  
(BMW, Bosch, Siemens, Boeing, John Deere, CAT, Panasonic, ...)

[INFO]

More info: <https://www.ros.org>



---

N e x T

# Hands-on exercises

Tools

# Google Colaboratory

- Write and execute Python in your browser
- Jupyter Notebook like
- Zero configuration required
- Access to AI-accelerators free of charge
  - TPU
  - Nvidia T4 GPU
- Pay for extra resources
  - Nvidia V100 GPU
  - Nvidia A100 GPU
  - ...
- Easy sharing
- Free:
  - Shared resources
  - Sessions are limited (time & resources)



A screenshot of a Google Colab notebook titled "CudaTestOnColab.ipynb". The code cell contains the command "!nvidia-smi". The output shows GPU information for a Tesla T4:

```
Mon Jan  8 08:49:11 2024
+-----+
| NVIDIA-SMI 535.104.05      Driver Version: 535.104.05    CUDA Version: 12.2 |
| Persistence-M | Bus-Id      Disp.A | Volatile Uncorr. ECC | | | | |
| GPU Name      Fan Temp   Perf  Pwr:Usage/Cap | Memory-Usage | GPU-Util Compute M. |
| Fan Temp   Perf |          |          |          |          |          |          |
+-----+
| 0  Tesla T4      Off  00000000:00:04.0 Off |           0% | Default |
| N/A 35C   P8      9W / 70W |    0MiB / 15360MiB |           0% |          |
+-----+
Processes:
| GPU  GI  CI      PID  Type  Process name          GPU Memory |
| ID   ID              |          |          |          | Usage     |
+-----+
| No running processes found
+-----+
```

[INFO]

More info: <https://colab.research.google.com>

# Libraries used in the hands-on exercises



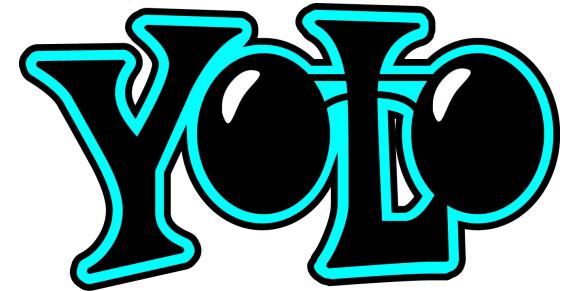
python<sup>TM</sup>

 pandas



seaborn

 NumPy

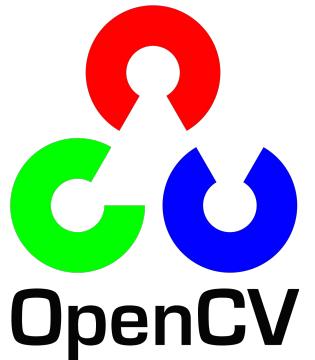
 YOP

 matplotlib

 Keras

 scikit-learn

 SciPy

 OpenCV



---

N e x T

## Hands-on exercises

<https://github.com/PXLAIRobotics/PG.AI-TA-2324-HandsOn>